How to deploy a NIM model in PC-AI – for beginners.

In this document I will show how to deploy a model in PC-AI using MLIS.

For this example, will use the NVIDIA-NIM vista 3D model (https://build.nvidia.com/nvidia/vista-3d)

Requirements:

Create an API key to be able to use the service, for NVIDIA NIM you need an NGC API key, for Hugging face you need a HF API key and so forth.

Overall steps:

During the process we will complete the following steps – I will comment on each step below.

0) **MLIS** – access MLIS within Private Cloud - AI
1) **Registries** - if there is no registry for NVIDIA-NIM models, you need to create one. In this example, I will create fra-onboarding
2) **Packaged models** – need to package the model. In this example, I will package NVIDIA-NIM vista 3D model under the name vista-3d-fra
3) **Deployments** – need to create a model deployment – In this example, the model wil be deployed under fra-onboarding name
4) **API-tokens** – need to generate an API token. This is needed to be able to make post requests to the model – in the example the api token will be called fra-onboarding.

Detailed step – by - step guide:

0) MLIS

- Go to **https://common.cloud.hpe.com/** and login.
- Choose a workspace
- Under **Featured Services** look for the **catalog** and find **Private Cloud AI** -> **Launch**
- In the dashboard you see a summary, the click on systems

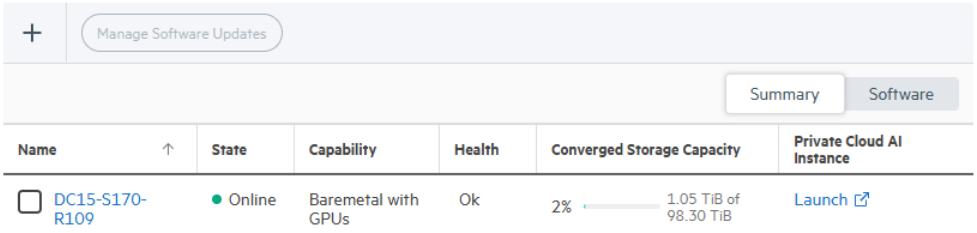

- Click on Launch.
- On the left look for tool and frameworks then under data science Open HPE MLIS

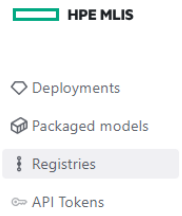

1) Registries

Example of a NGC registry (zoom the images when needed)

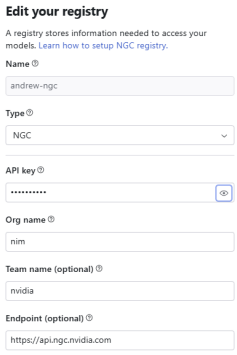

Example of an Hugging face registry (zoom the images when needed

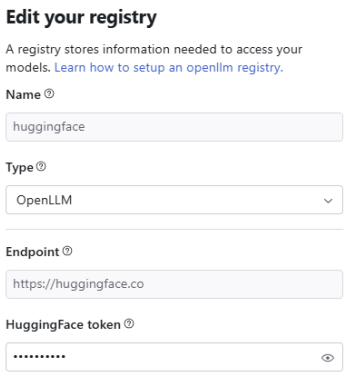

Create a registry

Click on `add new registry`

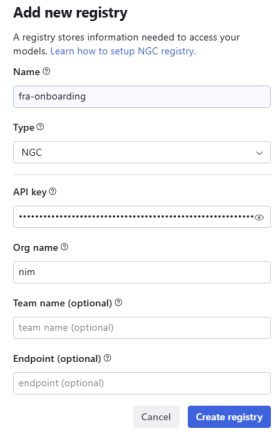

Name: your choice of name

Type: Since I am looking for models in NGC, In the registry type I chose NGC.

API key: The API KEY is your NGC API-KEY

Org name: I am looking for nim models, so I wrote **nim** as Org name (I noticed that if I used a different name I would not find nims model when I go to the next step – model package)

Team name nvidia and **Endpoint** since they are optional, they can be left empty and they will be automatically compiled.

Now that the registry is ready, I can package a model.

2) Packaged models

Click on **Add new model**

Your model: Chose the name and the description

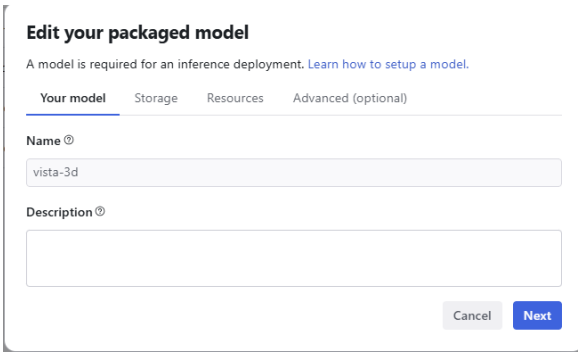

Storage: you should see the NIM model you are interested on in the drop down list. The image should self compile, **Path** is the field you can use if you'd like the model to be downloaded/cached once and stay in the persistent volume claim (PVC).

In this case if you type pvc://Kubeflow-shared-pvc/francesco it will create a folder there if it doesn't exist and the model will be cached there.

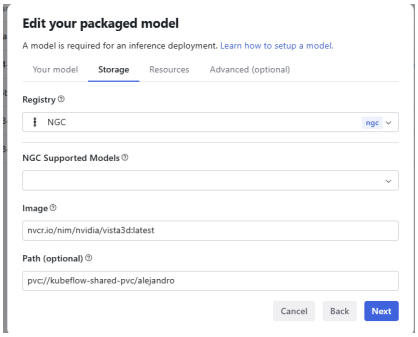

Note: If you don't see NVIDIA Vista among the models, you need to make sure that MLIS was deployed with disable_ngc set to false.

In AIE go to Tools and Frameworks, data science, MLIS config. At the very bottom

```
__internal:
  disable_ngc: false
```

Resources: This is an example for NVIDIA NIM VISTA3D, which is already available – zoom images as needed.

These are Alejandro's config for the model when it comes to deciding the resources

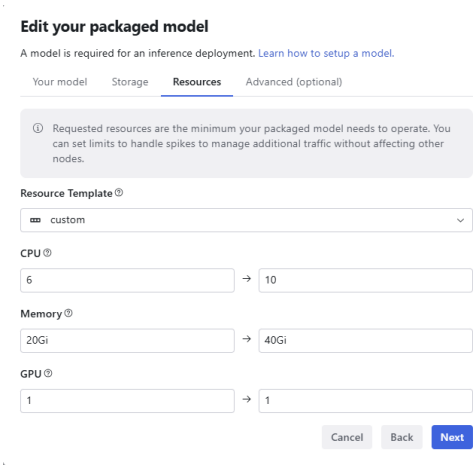

These are Andrew's:

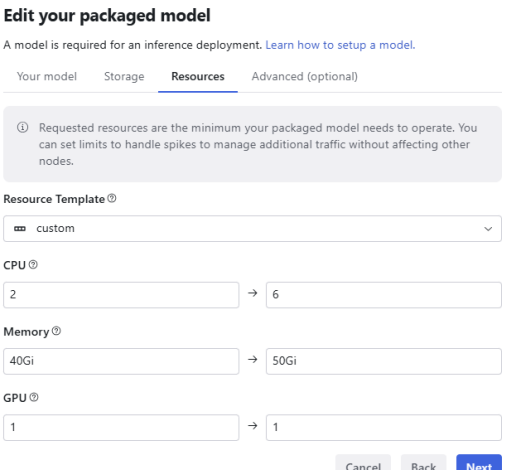

ADVANCED (optional): Alejandro left it empty

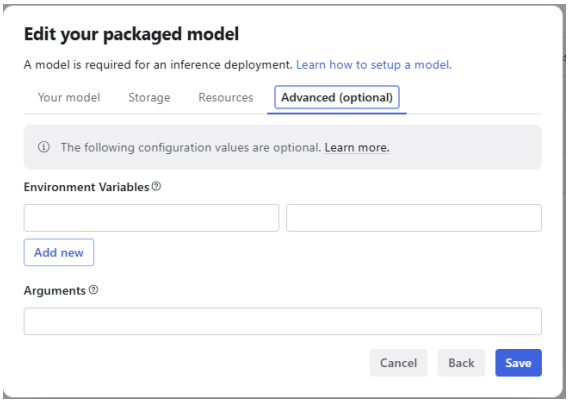

while Andrew didn't

And used variables from https://docs.nvidia.com/nim/medical/vista3d/latest/advanced-usage.html

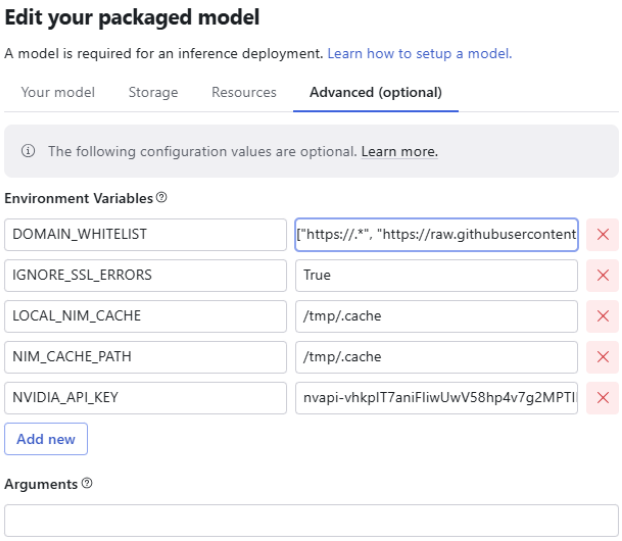

I noticed that if I left https://.*.ingress.pcai0109.dc15.hpecolo.net/.* out of the DOMAIN_WHITELIST hence, using the default values I was not able to make post requests to the model, and do inference, so I**decided to use the DOMAIN_WHITELIST like Andrew and included** left https://.*.ingress.pcai0109.dc15.hpecolo.net/.*

DOMAIN_WHITELIST = ["https://.*", "https://raw.githubusercontent.com/NVIDIA/.*", "https://assets.ngc.nvidia.com/products/api-catalog/vista3d/.*", "https://storage.googleapis.com/.*", "https://.*.s3.amazonaws.com/.*", "https://.*.blob.core.windows.net/.*",**"https://.*.ingress.pcai0109.dc15.hpecolo.net/.*"**]

Below are my settings:

Response:

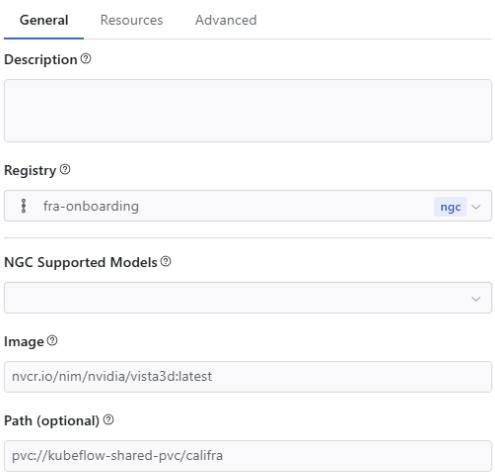

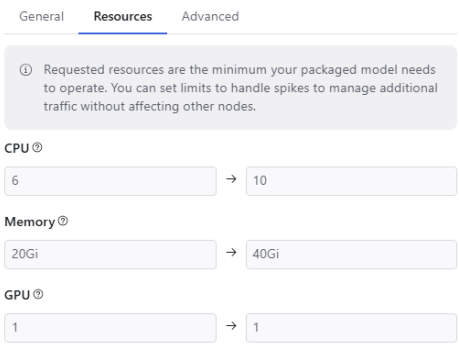

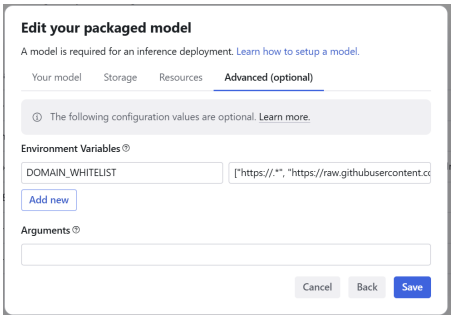

Now that the model is added it will appear as staged

Model name			Status	Last modified	Description	Registry used	Path
vista-3d-fra	v1	···	Staged	3 minutes ago		fra-onboarding	pvc://kubeflow-shared-pvc/califra

Deployments

Click on **create a new deployment**.

- **Deployment:** Choose a deployment name and description

- **Packaged Model:** Select the packaged model you'd like to deploy, and its version
- **Infrastructure:** Leave endpoint security on under infrastructure
- **Scaling:** Select the auto scaling target template. Here I selected the same that Andrew had used:
    - Autoscaling targets template **custom,**
    - minimum instance **1**
    - maximum instances **1**
    - auto scaling target **rps 0**

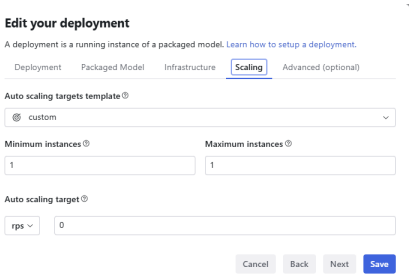

- **Advanced (optional) -** in this was left empty

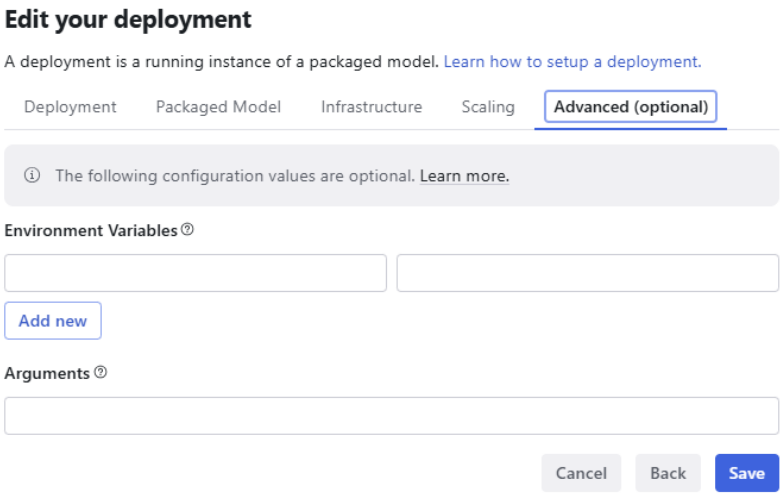

- **API Tokens**
  Create new API access token.

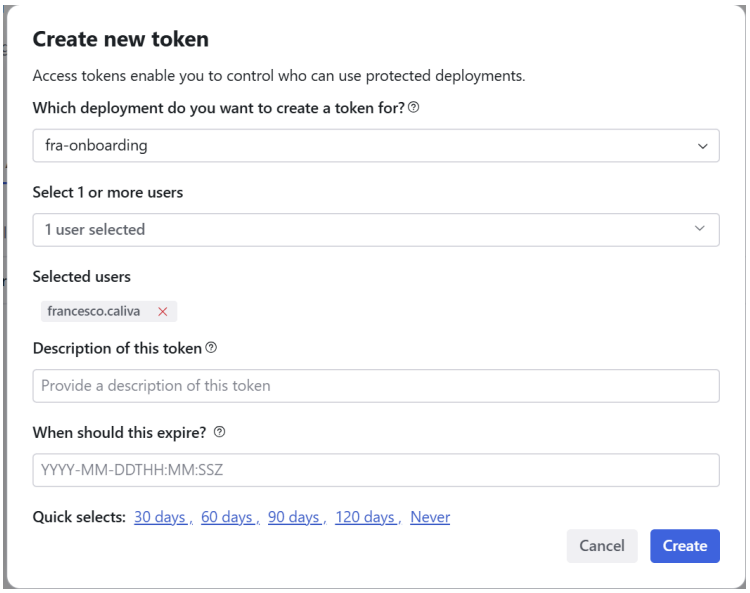

This API key is the key which you will use when sending post requests to the model for instance:

```
base_url="https://fra-onboarding-predictor-francesco-caliv-2a23f35c.ingress.pcai0109.dc15.hpecolo.net"

mlis_token = "eyJhbGciOiJSUzI1NiIsInR5cCI6IkpXVCJ9.eyJleHAiOjE3NTcyODQwMTcsImlhdCI6MTc1NDY5MjAxOCwiaXNzIjoiYWlvbGlAaHBlLmNvbSIsInN1YiI6ImQ3MGZiYjU4LTdmYmItNDcwNC1hZjRkLWFjZjAwMTYxYjhhMCIsInVzZXIiOiJhZG1pbiJ9.IuXF-gZ5UiROspdBKPw1XZyb-9mO-zPV13Cq6wnYjoKyu1ub5dVpeuthVVYQePRmaw8iV3sHAJkc3g3Dqx6jSkWTHZsGhlwnnKK5lBtNm0L2ApHQAFuD7sQvbFigJ3eGf2Mi3Sm8NcNIQDTvCiERvRbQXYe6S8JQ1GhfFv3I3cLU5xM8WnCtlBugRJeMp9_DPUfaZtdJj738FB0Pdnio8D19yVcHLXvOqO3ordX8enLQs8Wq3sJXNC5ZLVG5TtUVh_qEJL7y9EloPbMZwGUL2Zq8Ytodvxz9N3qjR8E0_utATJ38SMq_0ubC9nCC0juZGOLaaxQ9RtyfKE_BMCrAaA"

headers = {'Authorization': f'Bearer {mlis_token}'}"
response = requests.post(
        f'{base_url}/v1/vista3d/inference',
        json=data, headers=headers
)
```

```

Ideally it is recommended to add the api-key in a .env file and then load it to memory using dotenv.load_dotenv() function.

**Wait until it is ready and serving**


```